

Parallelizing Convex Learning Algorithms

Michael Curtis 260475694
Voula Theophanous 260480568

December 13, 2015

Abstract

This report outlines methodologies for parallelizing convex cost function machine learning algorithms. We present a Master-Worker implementation and a Gradient-Gossip based implementation. Both of these methods aim to minimize the amount of communications between processes and maximize the amount of parallel execution. We will also describe what a convex function is, how convexity relates to learning algorithms, and why convex learning algorithms can be paralyzed.

1 Introduction

The following example illustrates why executing a machine algorithms faster could be important. Imagine you are a computer scientist for the World Health Organization. There has been an outbreak of a mosquito born virus that has been growing in western African countries. It has been found that spraying airborne pesticides to control mosquito population is effective in controlling the spread of the virus. Your organization has 140 teams deligated to spraying the pesticides across the effected countries. You have over 3000 mosquito hotspots near population centers where mosquitos' blood have been testing at high rates for the virus. You are given the location of mosquito blood testing, associated rate of virus occurrence in mosquito and population centers, and previous pesticide sprays with associated location. You have decided to build a linear regression model aiming to minimize the total number of mosquito positive virus tests. Since there are hundreds or thousands of features in our training data the model takes days to build when running in serial. If we have a method to parallelize this learning algorithm we could send the 140 teams to an optimally chosen 140 of the 3000 locations in the magnitude of hours.

1.1 Definition: Convex Functions

The simple and informal definition of a convex function is to visualize a parabola, where we can easily see that there is only one global minimum. That is there is only one set of parameters to the derivative of the function that will return 0. Lets try to visualize this in general. Pick two sets of parameters to the function

that given us two different points on the functions. With the two points draw a line connecting, or in the higher dimensional case, visualize a surface where both the points are contained. If for any two points of any parameters has an obstruction disallowing a line of sight from one point to another, the function is not convex as there will be multiple parameter values such that the derivative of the function evaluates to 0. This is mathematically defined in Stanford Professor Stephen Boyd's textbook¹

A function $f : R^n \rightarrow R$ is *convex* if $domain f$ is a convex set and if for all $x, y \in domain f$, and θ with $0 \leq \theta \leq 1$, we have

$$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y) \quad (1)$$



Figure 1: A "chord" Between Two Points In a Convex Function

1.2 Convexity and Machine Learning

The objective of learning algorithms is to build a model using a set of training data X and target data y . The model can then take new X data and output a prediction of y data. In many cases describing the correlation between X and y is done by finding a weighting matrix w . As seen in figure 2², $w = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} 1.60 \\ 1.05 \end{bmatrix}$. These algorithms search for the optimal weight matrix by minimizing the error between X and y . Machine learning algorithms such as linear regression, support vector machines, LASSO linear regression, ridge regression, logistic regression, and metric learning all use **convex** cost functions to evaluate error.

Example 1.1. Linear regression builds a predictive model $Y = Xw$ by minimizing the least-squares, error evaluation function $error(w) = (Y - Xw)^T(Y - Xw)$. We minimize $error(w)$ by setting its derivative to zero and traversing potential w_0, w_1, \dots, w_n values until we are within an acceptable range of zero. Where we

¹https://web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf

²<http://cs.mcgill.ca/~jpineau/comp598/Lectures/02LinearRegression-publish.pdf>




Figure 2: w Describes The Line That Minimizes Error: $y = 1.60x + 1.05$

find a w such that $\frac{\partial \text{error}(w)}{\partial w} = -2X^T(Y - Xw) = 0$. This algorithm's output is illustrated in figure 2. Later in this report we will use linear regression to test our parallelization.

The key point of convexity as it relates to optimization of machine learning algorithm is that if a the error evaluation function is convex we can use a gradient descent algorithm to traverse the potential w_0, w_1, \dots, w_n values. Even more crucial is the fact that gradient descent algorithms are parallelizable.

Algorithm 1: Centralized Gradient Descent

```

given a starting point  $w \in \text{domain} f$ ;
repeat
    1.  $\Delta w := \nabla f(w)$ ;
    2. Line Search. Choose step size  $t$  via exact or backtracking line
       search.;
    3. Update.  $w := w + t\nabla f(w)$ 
until stopping criterion is satisfied;
```

We can visualize Algorithm 1's execution via the convex definition given in Figure 1 where the our gradient descent is traversing weight parameter values w_0, w_1, \dots, w_n . The algorithm stops once it has found the global minimum.

2 Master-Worker Parallelization

2.1 Implementation

Each evaluation of $\Delta w := \nabla f(w)$ in step 1 of Algorithm 1 can be computationally heavy with a large X . Our algorithm separates X into X_0, X_1, \dots, X_n and associates $\nabla f(w)$'s with n processes and thus, speeds up the time it takes for step 1 to execute on each iteration. Master-Worker delegates a single process to steps 2 and 3 of Algorithm 1, where each iteration's w vector is assigned, as well as the comparison epsilon check to see if the algorithm has satisfied the stopping criterion. This master process sends the newly found w to the n worker processes delegated to evaluating their local $\Delta w := \nabla f_i(w)$ then each worker sends that value back to the master for reduction via the summation $\nabla f(w) = \nabla f_0(w) + \nabla f_1(w) + \dots + \nabla f_n(w)$. Note that $\nabla f_i(w) = -2X_i^T(Y - X_i w)$ for $0 \leq i \leq n$ for our linear regression test example.

2.2 Results

3 Synchronous Gossip Parallelization

3.1 Implementation

Synchronous gossip is similar to the master-worker implementation in that the machine learning cost functions are required to be convex for parallelization. Synchronous gossip is fundamentally different because it leaves the responsibility of finding a weight matrix w_i to each process. In this section we will detail why all processes in the defined network would converge to the same solution for the weight matrix w . Algorithm 2 describes a process in which nodes compute gradients relative to their local gradient functions as defined by their X_i data.

Algorithm 2: Synchronous Gossip Gradient Descent

Neighbours at process i are $N_i = \{j : (i, j) \in Edges\} \cup \{i\}$;

Weighting matrix at process i is $x_0^i = [0, 0, \dots, 0]$;

repeat

1. Compute $\nabla f_i(x_k^i)$;
2. Send $\nabla f_i(x_k^i)$ to neighbors;
3. Receive $\nabla f_i(x_k^j)$ from neighbors;
4. Update $x_{k+1}^i = x_k^i - \alpha_k \sum_{j \in N_i} w_{ij} \nabla f_j(x_k^j)$

until *stopping criterion is satisfied*;

3.1.1 Convergence Requirements

Graph requirements for $G = (V, E)$:

1. G does not change

2. G is connected
3. G is undirected
4. $w_{ij} = 0$ if $(i, j) \notin E$ and $i \neq j$
5. $\sum_{j \in N_i} w_{ij} = 1$
6. $\sum_{i=1}^n w_{ij} = 1$

3.1.2 Communication Efficiency Requirements

There are two major contributing factors in our implementation that directly relate to an increased speedup of execution. The first of which is to decide whether there an edge exists between two processes by a probability that is a function of the number of processes. Overwhelming odds suggest that the graph created by this process, *Erdos-Renyi Model* equation (2), leads to a graph in which most nodes are connected by a edge distance of two. Having an average connection distance of two edges between two nodes allows for the most efficient communication of send and receives between processes.

$$Pr((i, j) \in E) = \sqrt{\frac{2\log(n)}{n}} \quad (2)$$




Figure 3: Graph Generated Via Erdos-Renyi Model

The second factor in allowing efficient communication in our synchronous gossip implementation is to assign edge weights in a way that leads to fast convergence. *Metropolis-Hastings*, equation 3, assign weights via edge degrees. In other words processes that have more influence over multiple nodes' weight

matrix value those processes receive more weighting to allow for faster communication of changes in the weights due to their local data.

$$\begin{aligned} \text{For } i \neq j \text{ set } w_{ij} &= \begin{cases} \frac{1}{\max(\deg(i), \deg(j))}, & \text{if } (i, j) \in E. \\ 0, & \text{otherwise.} \end{cases} \\ \text{For } i = j \text{ set } w_{ij} &= 1 - \sum_{j=1}^n w_{ij} \end{aligned} \tag{3}$$

3.2 Results

4 Conclusion

The key introductory knowledge is that some machine learning algorithms have convex cost functions, where convex cost functions allow for gradient descent and gradient descent is parallelizable. Therefore, we can generate substantial speedups.